

Introduction to Markov Decision Processes

INDE 597

March 6, 2019

Saumya Sinha

- Random variable X – function with many possible outcomes, each with some probability.

- Random variable X – function with many possible outcomes, each with some probability.
- Examples:
 - X = outcome of a (fair) coin toss: $\{H, T\}$.
 - Y = cars sold in Houston till date in 2019: $\{0, 1, \dots, M_1\}$.
 - Z = length of queue at the Coffeehouse at noon: $\{0, 1, \dots, M_2\}$.

- Random variable X – function with many possible outcomes, each with some probability.
- Examples:
 - X = outcome of a (fair) coin toss: $\{H, T\}$.
 - Y = cars sold in Houston till date in 2019: $\{0, 1, \dots, M_1\}$.
 - Z = length of queue at the Coffeehouse at noon: $\{0, 1, \dots, M_2\}$.
- Stochastic process $\{X_t\}_{t=1}^T$ – family of random variables, usually indexed by time.

- Random variable X – function with many possible outcomes, each with some probability.
- Examples:
 - X = outcome of a (fair) coin toss: $\{H, T\}$.
 - Y = cars sold in Houston till date in 2019: $\{0, 1, \dots, M_1\}$.
 - Z = length of queue at the Coffeehouse at noon: $\{0, 1, \dots, M_2\}$.
- Stochastic process $\{X_t\}_{t=1}^T$ – family of random variables, usually indexed by time.
- Examples:
 - (X_1, \dots, X_{20}) : outcomes of 20 (fair) coin tosses.
 - (Y_1, \dots, Y_{365}) : cars sold in Houston up to day t of 2019.
 - (Z_1, \dots, Z_T) : length of queue at the Coffeehouse every hour.

Markov Processes

- Stochastic process is 'Markovian' if X_{t+1} depends only on X_t (and not on X_1, \dots, X_{t-1}).

Formally, $\mathbb{P}(X_{t+1}|X_t) = \mathbb{P}(X_{t+1}|X_1, \dots, X_t)$.

Markov Processes

- Stochastic process is 'Markovian' if X_{t+1} depends only on X_t (and not on X_1, \dots, X_{t-1}).

Formally, $\mathbb{P}(X_{t+1}|X_t) = \mathbb{P}(X_{t+1}|X_1, \dots, X_t)$.

- Examples:

- (X_1, \dots, X_{20}) : outcomes of 20 (fair) coin tosses

Each toss independent of all previous tosses, hence Markovian.

Markov Processes

- Stochastic process is 'Markovian' if X_{t+1} depends only on X_t (and not on X_1, \dots, X_{t-1}).

Formally, $\mathbb{P}(X_{t+1}|X_t) = \mathbb{P}(X_{t+1}|X_1, \dots, X_t)$.

- Examples:

- (X_1, \dots, X_{20}) : outcomes of 20 (fair) coin tosses

Each toss independent of all previous tosses, hence Markovian.

- (Y_1, \dots, Y_{365}) : cars sold in Houston up to day t of 2019.

$Y_{t+1} = Y_t + \text{number of cars sold on day } t$, Markovian.

Markov Processes

- Stochastic process is ‘Markovian’ if X_{t+1} depends only on X_t (and not on X_1, \dots, X_{t-1}).

Formally, $\mathbb{P}(X_{t+1}|X_t) = \mathbb{P}(X_{t+1}|X_1, \dots, X_t)$.

- Examples:

- (X_1, \dots, X_{20}) : outcomes of 20 (fair) coin tosses

Each toss independent of all previous tosses, hence Markovian.

- (Y_1, \dots, Y_{365}) : cars sold in Houston up to day t of 2019.

$Y_{t+1} = Y_t +$ number of cars sold on day t , Markovian.

- (Z_1, \dots, Z_T) : length of queue at the Coffeehouse every hour.

$Z_{t+1} = Z_t +$ number of arrivals in hour t – number of people served in hour t , Markovian.

Markov Processes

- Stochastic process is ‘Markovian’ if X_{t+1} depends only on X_t (and not on X_1, \dots, X_{t-1}).

Formally, $\mathbb{P}(X_{t+1}|X_t) = \mathbb{P}(X_{t+1}|X_1, \dots, X_t)$.

- Examples:

- (X_1, \dots, X_{20}) : outcomes of 20 (fair) coin tosses

Each toss independent of all previous tosses, hence Markovian.

- (Y_1, \dots, Y_{365}) : cars sold in Houston up to day t of 2019.

$Y_{t+1} = Y_t + \text{number of cars sold on day } t$, Markovian.

- (Z_1, \dots, Z_T) : length of queue at the Coffeehouse every hour.

$Z_{t+1} = Z_t + \text{number of arrivals in hour } t - \text{number of people served in hour } t$, Markovian.

- Bag with 5 red and 5 black balls, $W_t = \text{color of } t\text{-th ball drawn}$, $t = 1, \dots, 10$.

If balls are replaced, all draws are independent; Markov process.

Markov Processes

- Stochastic process is ‘Markovian’ if X_{t+1} depends only on X_t (and not on X_1, \dots, X_{t-1}).

Formally, $\mathbb{P}(X_{t+1}|X_t) = \mathbb{P}(X_{t+1}|X_1, \dots, X_t)$.

- Examples:

- (X_1, \dots, X_{20}) : outcomes of 20 (fair) coin tosses

Each toss independent of all previous tosses, hence Markovian.

- (Y_1, \dots, Y_{365}) : cars sold in Houston up to day t of 2019.

$Y_{t+1} = Y_t + \text{number of cars sold on day } t$, Markovian.

- (Z_1, \dots, Z_T) : length of queue at the Coffeehouse every hour.

$Z_{t+1} = Z_t + \text{number of arrivals in hour } t - \text{number of people served in hour } t$, Markovian.

- Bag with 5 red and 5 black balls, $W_t = \text{color of } t\text{-th ball drawn}$, $t = 1, \dots, 10$.

If balls are replaced, all draws are independent; Markov process.

If balls are not replaced, W_{t+1} depends on all previous draws; not Markovian.

- Stochastic process $\{s_t\}$, $t = 0, 1, \dots, T$.

Markov Decision Processes

- Stochastic process $\{s_t\}$, $t = 0, 1, \dots, T$.
- Every t , an agent observes s_t , and takes an action a_t .

Markov Decision Processes

- Stochastic process $\{s_t\}$, $t = 0, 1, \dots, T$.
- Every t , an agent observes s_t , and takes an action a_t .
- Next state s_{t+1} depends on s_t (Markov) and a_t (decision).

Markov Decision Processes

- Stochastic process $\{s_t\}$, $t = 0, 1, \dots, T$.
- Every t , an agent observes s_t , and takes an action a_t .
- Next state s_{t+1} depends on s_t (Markov) and a_t (decision).
- Example: (Z_1, \dots, Z_T) : length of queue at the Coffeehouse every hour.

Markov Decision Processes

- Stochastic process $\{s_t\}$, $t = 0, 1, \dots, T$.
- Every t , an agent observes s_t , and takes an action a_t .
- Next state s_{t+1} depends on s_t (Markov) and a_t (decision).
- Example: (Z_1, \dots, Z_T) : length of queue at the Coffeehouse every hour.
 - Suppose each counter serves 30 people in an hour.

Markov Decision Processes

- Stochastic process $\{s_t\}$, $t = 0, 1, \dots, T$.
- Every t , an agent observes s_t , and takes an action a_t .
- Next state s_{t+1} depends on s_t (Markov) and a_t (decision).
- Example: (Z_1, \dots, Z_T) : length of queue at the Coffeehouse every hour.
 - Suppose each counter serves 30 people in an hour.
 - At the start of hour t :

Markov Decision Processes

- Stochastic process $\{s_t\}$, $t = 0, 1, \dots, T$.
- Every t , an agent observes s_t , and takes an action a_t .
- Next state s_{t+1} depends on s_t (Markov) and a_t (decision).
- Example: (Z_1, \dots, Z_T) : length of queue at the Coffeehouse every hour.
 - Suppose each counter serves 30 people in an hour.
 - At the start of hour t :
 - observe Z_t .

Markov Decision Processes

- Stochastic process $\{s_t\}$, $t = 0, 1, \dots, T$.
- Every t , an agent observes s_t , and takes an action a_t .
- Next state s_{t+1} depends on s_t (Markov) and a_t (decision).
- Example: (Z_1, \dots, Z_T) : length of queue at the Coffeehouse every hour.
 - Suppose each counter serves 30 people in an hour.
 - At the start of hour t :
 - observe Z_t .
 - decide how many counters to open (say $a_t = 1, 2$ or 3).

Markov Decision Processes

- Stochastic process $\{s_t\}$, $t = 0, 1, \dots, T$.
- Every t , an agent observes s_t , and takes an action a_t .
- Next state s_{t+1} depends on s_t (Markov) and a_t (decision).
- Example: (Z_1, \dots, Z_T) : length of queue at the Coffeehouse every hour.
 - Suppose each counter serves 30 people in an hour.
 - At the start of hour t :
 - observe Z_t .
 - decide how many counters to open (say $a_t = 1, 2$ or 3).
 - then, $Z_{t+1} = Z_t - 30a_t + \text{number of new arrivals}$.

Markov Decision Processes

- Stochastic process $\{s_t\}$, $t = 0, 1, \dots, T$.
- Every t , an agent observes s_t , and takes an action a_t .
- Next state s_{t+1} depends on s_t (Markov) and a_t (decision).
- Example: (Z_1, \dots, Z_T) : length of queue at the Coffeehouse every hour.
 - Suppose each counter serves 30 people in an hour.
 - At the start of hour t :
 - observe Z_t .
 - decide how many counters to open (say $a_t = 1, 2$ or 3).
 - then, $Z_{t+1} = Z_t - 30a_t + \text{number of new arrivals}$.
 - profit = revenue from coffees sold - wages paid to servers.

Markov Decision Processes

- Stochastic process $\{s_t\}$, $t = 0, 1, \dots, T$.
- Every t , an agent observes s_t , and takes an action a_t .
- Next state s_{t+1} depends on s_t (Markov) and a_t (decision).
- Example: (Z_1, \dots, Z_T) : length of queue at the Coffeehouse every hour.
 - Suppose each counter serves 30 people in an hour.
 - At the start of hour t :
 - observe Z_t .
 - decide how many counters to open (say $a_t = 1, 2$ or 3).
 - then, $Z_{t+1} = Z_t - 30a_t + \text{number of new arrivals}$.
 - profit = revenue from coffees sold - wages paid to servers.
 - Goal: maximize total profit.

Markov Decision Processes

- A Markov decision process (MDP) consists of:
 - Decision epochs $t = 0, 1, \dots, T - 1$.
 - State space $\mathcal{S} = \{1, 2, \dots, S\}$.
 - Action space $\mathcal{A} = \{1, 2, \dots, A\}$.
 - Transition probabilities $P_a(s, s') = \mathbb{P}(s_{t+1} = s' | s_t = s, a_t = a)$.
 - Rewards $R(s, a)$.

Markov Decision Processes

- A Markov decision process (MDP) consists of:
 - Decision epochs $t = 0, 1, \dots, T - 1$.
 - State space $\mathcal{S} = \{1, 2, \dots, S\}$.
 - Action space $\mathcal{A} = \{1, 2, \dots, A\}$.
 - Transition probabilities $P_a(s, s') = \mathbb{P}(s_{t+1} = s' | s_t = s, a_t = a)$.
 - Rewards $R(s, a)$.
- Policy/decision rule π assigns an action to every state.

Markov Decision Processes

- A Markov decision process (MDP) consists of:
 - Decision epochs $t = 0, 1, \dots, T - 1$.
 - State space $\mathcal{S} = \{1, 2, \dots, S\}$.
 - Action space $\mathcal{A} = \{1, 2, \dots, A\}$.
 - Transition probabilities $P_a(s, s') = \mathbb{P}(s_{t+1} = s' | s_t = s, a_t = a)$.
 - Rewards $R(s, a)$.
- Policy/decision rule π assigns an action to every state.
- Total expected reward for a policy

$$V^\pi(s_0) = \mathbb{E} \left[\sum_{t=0}^{T-1} R(s_t, \pi(s_t)) + r(s_T) \right],$$

where r is a terminal reward.

Markov Decision Processes

- A Markov decision process (MDP) consists of:
 - Decision epochs $t = 0, 1, \dots, T - 1$.
 - State space $\mathcal{S} = \{1, 2, \dots, S\}$.
 - Action space $\mathcal{A} = \{1, 2, \dots, A\}$.
 - Transition probabilities $P_a(s, s') = \mathbb{P}(s_{t+1} = s' | s_t = s, a_t = a)$.
 - Rewards $R(s, a)$.
- Policy/decision rule π assigns an action to every state.
- Total expected reward for a policy

$$V^\pi(s_0) = \mathbf{E} \left[\sum_{t=0}^{T-1} R(s_t, \pi(s_t)) + r(s_T) \right],$$

where r is a terminal reward.

- V^π is called the value (function) for policy π .

Markov Decision Processes

- A Markov decision process (MDP) consists of:
 - Decision epochs $t = 0, 1, \dots, T - 1$.
 - State space $\mathcal{S} = \{1, 2, \dots, S\}$.
 - Action space $\mathcal{A} = \{1, 2, \dots, A\}$.
 - Transition probabilities $P_a(s, s') = \mathbb{P}(s_{t+1} = s' | s_t = s, a_t = a)$.
 - Rewards $R(s, a)$.
- Policy/decision rule π assigns an action to every state.
- Total expected reward for a policy

$$V^\pi(s_0) = \mathbb{E} \left[\sum_{t=0}^{T-1} R(s_t, \pi(s_t)) + r(s_T) \right],$$

where r is a terminal reward.

- V^π is called the value (function) for policy π .
- Goal: find π that maximizes this reward. Optimal value function

$$V^*(s_0) = \max_{\pi} V^\pi(s_0).$$

- Determine optimal service schedule for incoming jobs.

Examples: Queuing

- Determine optimal service schedule for incoming jobs.
- Decision epochs – every hour.
- State = number of jobs waiting to be served.
- Action = number of jobs accepted for service in each period, $\mathcal{A} = 0, 1, \dots, s_t$.
- Next state = $s_{t+1} = s_t - a_t + \text{new arrivals (random)}$.
- Reward for completing a job, cost for keeping jobs waiting.

Examples: Inventory management

- Determine optimal policy for ordering inventory in the face of uncertain demand.

Examples: Inventory management

- Determine optimal policy for ordering inventory in the face of uncertain demand.
- Decision epochs – every week.
- Seller observes the current inventory (state) and decides how much more inventory (action) to acquire for the week.
- Random demand d_t during the week.
- $s_{t+1} = \max\{s_t + a_t - d_t, 0\}$.
- Reward = sale revenue - purchase cost.

Example: medical treatment planning

- Find optimal dosing policy for 'best' disease progression.
- Observe a 'disease score' (state) of a patient, and prescribe a dose (action) of medication (rheumatoid arthritis, radiation therapy).
- $s_{t+1} = f(s_t, a_t)$ clinically determined.
- Reward = improvement in patient's condition.

Solving MDPs: Bellman's Principle of Optimality

Recall objective

$$V^*(s_0) = \max_{\pi} V^{\pi}(s_0) = \max_{\pi} \mathbb{E} \left[\sum_{t=0}^{T-1} R(s_t, \pi(s_t)) + r(s_T) \right].$$

Solving MDPs: Bellman's Principle of Optimality

Recall objective

$$V^*(s_0) = \max_{\pi} V^{\pi}(s_0) = \max_{\pi} \mathbb{E} \left[\sum_{t=0}^{T-1} R(s_t, \pi(s_t)) + r(s_T) \right].$$

Define 'rewards-to-go':

$$V_t^{\pi}(s_t) = \mathbb{E} \left[\sum_{t=t}^{T-1} R(s_t, \pi(s_t)) + r(s_T) \right]$$

$$V_t^*(s_t) = \max_{\pi} \mathbb{E} \left[\sum_{t=t}^{T-1} R(s_t, \pi(s_t)) + r(s_T) \right],$$

Solving MDPs: Bellman's Principle of Optimality

Recall objective

$$V^*(s_0) = \max_{\pi} V^{\pi}(s_0) = \max_{\pi} \mathbb{E} \left[\sum_{t=0}^{T-1} R(s_t, \pi(s_t)) + r(s_T) \right].$$

Define 'rewards-to-go':

$$V_t^{\pi}(s_t) = \mathbb{E} \left[\sum_{t=t}^{T-1} R(s_t, \pi(s_t)) + r(s_T) \right]$$

$$V_t^*(s_t) = \max_{\pi} \mathbb{E} \left[\sum_{t=t}^{T-1} R(s_t, \pi(s_t)) + r(s_T) \right],$$

Bellman's principle of optimality

$$V_t^*(s_t) = \max_{a \in \mathcal{A}} \left\{ R(s_t, a) + \lambda V_{t+1}^*(s_{t+1}) \right\}$$

$$V_T^*(s_T) = r(s_T).$$

Solving MDPs: Bellman's Principle of Optimality

Recall objective

$$V^*(s_0) = \max_{\pi} V^{\pi}(s_0) = \max_{\pi} \mathbb{E} \left[\sum_{t=0}^{T-1} R(s_t, \pi(s_t)) + r(s_T) \right].$$

Define 'rewards-to-go':

$$V_t^{\pi}(s_t) = \mathbb{E} \left[\sum_{t=t}^{T-1} R(s_t, \pi(s_t)) + r(s_T) \right]$$

$$V_t^*(s_t) = \max_{\pi} \mathbb{E} \left[\sum_{t=t}^{T-1} R(s_t, \pi(s_t)) + r(s_T) \right],$$


Bellman's principle of optimality

$$V_t^*(s_t) = \max_{a \in \mathcal{A}} \left\{ R(s_t, a) + \lambda V_{t+1}^*(s_{t+1}) \right\}$$

$$V_T^*(s_T) = r(s_T).$$

Bellman equations, solved by backward recursion.

Solving MDPs: Bellman's Principle of Optimality



$$V^*(s_0) = \max \left\{ R(s_0, a^1) + \lambda V_1^*(s_1^1), R(s_0, a^2) + \lambda V_1^*(s_1^2) \right\}.$$

Infinite-horizon MDPs

- Decision epochs $t = 0, 1, \dots$
- State space $\mathcal{S} = \{1, 2, \dots, S\}$.
- Action space $\mathcal{A} = \{1, 2, \dots, A\}$.
- Transition probabilities $P_a(s, s') = \mathbb{P}(s_{t+1} = s' | s_t = s, a_t = a)$.
- Rewards $R(s, a)$.

Infinite-horizon MDPs

- Decision epochs $t = 0, 1, \dots$
- State space $\mathcal{S} = \{1, 2, \dots, S\}$.
- Action space $\mathcal{A} = \{1, 2, \dots, A\}$.
- Transition probabilities $P_a(s, s') = \mathbb{P}(s_{t+1} = s' | s_t = s, a_t = a)$.
- Rewards $R(s, a)$.

Total expected reward for a policy

$$V^\pi(s_0) = \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t R(s_t, \pi(s_t)) \right],$$

where $\lambda \in (0, 1)$ is a discount factor.

Infinite-horizon MDPs

- Decision epochs $t = 0, 1, \dots$
- State space $\mathcal{S} = \{1, 2, \dots, S\}$.
- Action space $\mathcal{A} = \{1, 2, \dots, A\}$.
- Transition probabilities $P_a(s, s') = \mathbb{P}(s_{t+1} = s' | s_t = s, a_t = a)$.
- Rewards $R(s, a)$.

Total expected reward for a policy

$$V^\pi(s_0) = \mathbb{E} \left[\sum_{t=0}^{\infty} \lambda^t R(s_t, \pi(s_t)) \right],$$

where $\lambda \in (0, 1)$ is a discount factor.

No 'terminal' reward.

Solving MDPs: Bellman equations

- Principle of optimality still holds.

Solving MDPs: Bellman equations

- Principle of optimality still holds.
- Bellman equations

$$V^*(s) = \max_{a \in \mathcal{A}} \left\{ R(s, a) + \lambda \mathbb{E}[V^*(s')] \right\}.$$

Solving MDPs: Bellman equations

- Principle of optimality still holds.
- Bellman equations

$$V^*(s) = \max_{a \in \mathcal{A}} \left\{ R(s, a) + \lambda \mathbb{E}[V^*(s')] \right\}.$$

- No backward recursion.

Solving MDPs: Bellman equations

- Principle of optimality still holds.
- Bellman equations

$$V^*(s) = \max_{a \in \mathcal{A}} \left\{ R(s, a) + \lambda \mathbb{E}[V^*(s')] \right\}.$$

- No backward recursion.
- Solution methods:
 - Value iteration
 - Policy iteration
 - Linear programming

- Define Bellman operator

$$\mathcal{L}(v)(s) = \max_{a \in \mathcal{A}} R(s, a) + \lambda \mathbb{E}[v(s')].$$

- V^* is the fixed point of this operator.
- Value iteration relies on the Contraction Mapping Theorem.
- Algorithm:
 - Initial guess $V^{(0)}$.
 - Compute $V^{(k)} = \mathcal{L}(V^{(k-1)})$.
 - Iterate until stopping condition.

- Now we iterate over policies to find a better one.
- Algorithm:
 - Initial guess $\pi^{(0)}$.
 - Evaluate policy using matrix inversion:

$$V^{\pi^{(k)}}(s) = R(s, \pi(s)) + \mathbf{E}[V^{\pi^{(k)}}(s')].$$

- Improve policy

$$\pi^{(k+1)}(s) = \operatorname{argmax}_{a \in \mathcal{A}} \left\{ R(s, a) + \mathbf{E}[V^{\pi^{(k)}}(s')] \right\}.$$

- Iterate until stopping condition.

- Note that

$$\begin{aligned} V^*(s) &= \max_{a \in \mathcal{A}} \left\{ R(s, a) + \mathbf{E}[V^*(s')] \right\} \quad \forall s \\ &\geq R(s, a) + \mathbf{E}[V^*(s')] \quad \forall s, a. \end{aligned}$$

- V^* is the solution to the following LP:

$$\begin{aligned} \min \quad & \sum_{s \in \mathcal{S}} \beta(s) v(s) \\ \text{s.t.} \quad & v(s) \geq R(s, a) + \mathbf{E}[v(s')] \quad \forall s \in \mathcal{S}, a \in \mathcal{A}. \end{aligned}$$

Conclusion

- MDPs form a flexible modeling environment for sequential decision making problems.
- Several extensions – partially observable MDPs, reinforcement learning.
- Limitations:
 - ‘Curse of dimensionality’
 - Data requirements.